

WorldFork Overnight Accuracy Evaluation

Local Codex Harness

2026-04-30 13:04:39 PDT

Scope

This local-only study evaluates WorldFork against 50 anonymized concluded-event dossiers arranged as 10 categories by 5 horizons. Hidden source and outcome files stay separate from prompt dossiers. All live API-credit calls are constrained to `google/gemini-3.1-flash-lite-preview`.

Primary Reproducibility Reference

The running log is `agent-testing/accuracy-overnight/REPRODUCIBILITY_LOG.md`. Generated run records, command outputs, source files, event prompts, score tables, and this $\text{\LaTeX}$ source are saved under `agent-testing/accuracy-overnight/`.

Aggregate Results

Metric	Value
Completed runs	109
Top-outcome match rate	0.8073
Mean total variation distance	0.3639
Mean actual-outcome probability	0.4417
Gemini-only completed runs	109
Audited LLM calls	6439

Reliability Pattern

Baseline completed 50/50 runs. The sweep completed 59/64 runs. All five incomplete sweep cells were `long_high_detail_strict` and failed after late LLM `unavailable` tick errors left one or more active multiverses, causing the final report endpoint to reject nonterminal inputs. A direct final-report retry repeats the precondition failure until those multiverses are advanced or explicitly made reportable. The local harness now includes `resume-failures`, fresh idempotency keys, transient backoff, preserved `resume_*` command records, model-audit refresh, and a fail-fast `transient_failure_cap_reached` status for provider outages.

Incomplete Run Details

Event	Status	Failed Tick	Active Multiverses
<code>labor_3y_low_wage_campaign</code>	<code>report_failed</code>	<code>tick_008</code>	M1, M1.1

labor_1_3y_coffee_union_campaign	report_failed	tick_008	M1
tech_3y_flesharing_injunction	report_failed	tick_010	M1, M1.1
ai_1_3y_debt_recovery_scandal	transient_failure_cap_reached	tick_008	M1
ai_3y_childcare_benefit_scandal	report_failed	tick_009	M1, M1.1

Variant Breakdown

Variant	Count	Match	Mean TVD	Actual Prob.
baseline	50	0.78	0.3536	0.4434
long_compressed_default	16	0.75	0.3781	0.4385
long_high_detail_strict	11	0.9091	0.3739	0.3833
short_compressed_permissive	16	0.8125	0.3844	0.4677
short_high_detail_default	16	0.875	0.3549	0.4539

Category Breakdown

Category	Count	Match	Mean TVD	Actual Prob.
AI/public systems	23	0.8261	0.3561	0.4312
campus/civil society	5	0.8	0.3133	0.3617
corporate PR crisis	5	1.0	0.3589	0.5571
elections/legitimacy	9	0.4444	0.3744	0.3611
environment/resource	5	1.0	0.2747	0.4867
finance/market confidence	5	0.8	0.3893	0.5067
labor/social movements	23	0.8261	0.3817	0.3962
policy/regulatory backlash	5	0.8	0.268	0.4614
public health	5	0.4	0.424	0.3
tech/platform	24	0.9167	0.3829	0.5208

Horizon Breakdown

Horizon	Count	Match	Mean TVD	Actual Prob.
1-3 months	22	0.8636	0.3283	0.3848
1-3 years	20	0.7	0.4285	0.4112
3+ years	19	1.0	0.3046	0.4977
3-12 months	22	0.9091	0.3181	0.5562
weeks	26	0.6154	0.4266	0.3755

Findings and Advice

The strongest product direction is to make endpoint assessment explicit as durable state, not additional prompt wording. Endpoint guidance and outcome conclusions already exist in prompts/reports; the novel gaps are endpoint coverage ledgers, process-vs-terminal state gates, authority-weighted endpoint

evidence, contradiction checks over live alternatives, and outcome-label-aware scoring. For robust evaluation, `short_high_detail_default` was the best completed sweep setting by reliability-adjusted signal, while `long_high_detail_strict` should be reserved for smaller pilots with retry and branch-drain safeguards.

Reproducibility

The benchmark catalog, hidden source files, raw command outputs, model audits, score summaries, Markdown reports, this $\text{\LaTeX}$ source, and the compiled PDF are retained locally. No generated file is deleted by this harness.